wjafkibqs

July 31, 2023

[12]:
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

[13]:
```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
```

[14]:
```
df=pd.read_csv("/content/drive/MyDrive/mydatasets/21_cities.csv")
df
```

[14]:

	id	name	state_id	state_code	state_name	\
0	52	Ashkāsham	3901	BDS	Badakhshan	
1	68	Fayzabad	3901	BDS	Badakhshan	
2	78	Jurm	3901	BDS	Badakhshan	
3	84	Khandūd	3901	BDS	Badakhshan	
4	115	Rāghistān	3901	BDS	Badakhshan	
…	…	…	…	…	…	
150449	131496	Redcliff	1957	MI	Midlands Province	
150450	131502	Shangani	1957	MI	Midlands Province	
150451	131503	Shurugwi	1957	MI	Midlands Province	
150452	131504	Shurugwi District	1957	MI	Midlands Province	
150453	131508	Zvishavane District	1957	MI	Midlands Province	

	country_id	country_code	country_name	latitude	longitude	wikiDataId
0	1	AF	Afghanistan	36.68333	71.53333	Q4805192
1	1	AF	Afghanistan	37.11664	70.58002	Q156558
2	1	AF	Afghanistan	36.86477	70.83421	Q10308323
3	1	AF	Afghanistan	36.95127	72.31800	Q3290334
4	1	AF	Afghanistan	37.66079	70.67346	Q2670909
…	…	…	…	…	…	…
150449	247	ZW	Zimbabwe	-19.03333	29.78333	Q584001
150450	247	ZW	Zimbabwe	-19.78333	29.36667	Q32017959
150451	247	ZW	Zimbabwe	-19.67016	30.00589	Q32019023
150452	247	ZW	Zimbabwe	-19.75000	30.16667	Q7505444

```
150453         247           ZW     Zimbabwe -20.30345   30.07514  Q24235929

[150454 rows x 11 columns]
```

[15]:
```
df.head()
```

[15]:
```
     id       name  state_id state_code  state_name  country_id country_code  \
0    52  Ashkāsham      3901        BDS  Badakhshan           1           AF
1    68   Fayzabad      3901        BDS  Badakhshan           1           AF
2    78       Jurm      3901        BDS  Badakhshan           1           AF
3    84    Khandūd      3901        BDS  Badakhshan           1           AF
4   115  Rāghistān      3901        BDS  Badakhshan           1           AF

  country_name  latitude  longitude wikiDataId
0  Afghanistan  36.68333   71.53333   Q4805192
1  Afghanistan  37.11664   70.58002    Q156558
2  Afghanistan  36.86477   70.83421  Q10308323
3  Afghanistan  36.95127   72.31800   Q3290334
4  Afghanistan  37.66079   70.67346   Q2670909
```

1 Data Cleaning and Data Preprocessing

[16]:
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150454 entries, 0 to 150453
Data columns (total 11 columns):
 #   Column        Non-Null Count   Dtype
---  ------        --------------   -----
 0   id            150454 non-null  int64
 1   name          150454 non-null  object
 2   state_id      150454 non-null  int64
 3   state_code    150129 non-null  object
 4   state_name    150454 non-null  object
 5   country_id    150454 non-null  int64
 6   country_code  150406 non-null  object
 7   country_name  150454 non-null  object
 8   latitude      150454 non-null  float64
 9   longitude     150454 non-null  float64
 10  wikiDataId    147198 non-null  object
dtypes: float64(2), int64(3), object(6)
memory usage: 12.6+ MB
```

[17]:
```
df.describe()
```

[17]:
```
                 id       state_id     country_id       latitude  \
count  150454.000000  150454.000000  150454.000000  150454.000000
mean    76407.091689    2678.377677     140.658460      31.556175
std     44357.755335    1363.513591      70.666123      22.813220
min         1.000000       1.000000       1.000000     -75.000000
25%     38160.250000    1451.000000      82.000000      19.000000
50%     75975.500000    2174.000000     142.000000      40.684720
75%    115204.750000    3905.000000     207.000000      47.239220
max    153528.000000    5116.000000     247.000000      73.508190

           longitude
count  150454.000000
mean        2.369557
std        68.012770
min      -179.121980
25%       -58.468150
50%         8.669980
75%        27.750000
max       179.466000
```

[18]:
```
df.columns
```

[18]:
```
Index(['id', 'name', 'state_id', 'state_code', 'state_name', 'country_id',
       'country_code', 'country_name', 'latitude', 'longitude', 'wikiDataId'],
      dtype='object')
```

2 EDA and Visualization

[19]:
```
sns.pairplot(df)
```

[19]:
```
<seaborn.axisgrid.PairGrid at 0x7eea24d36050>
```

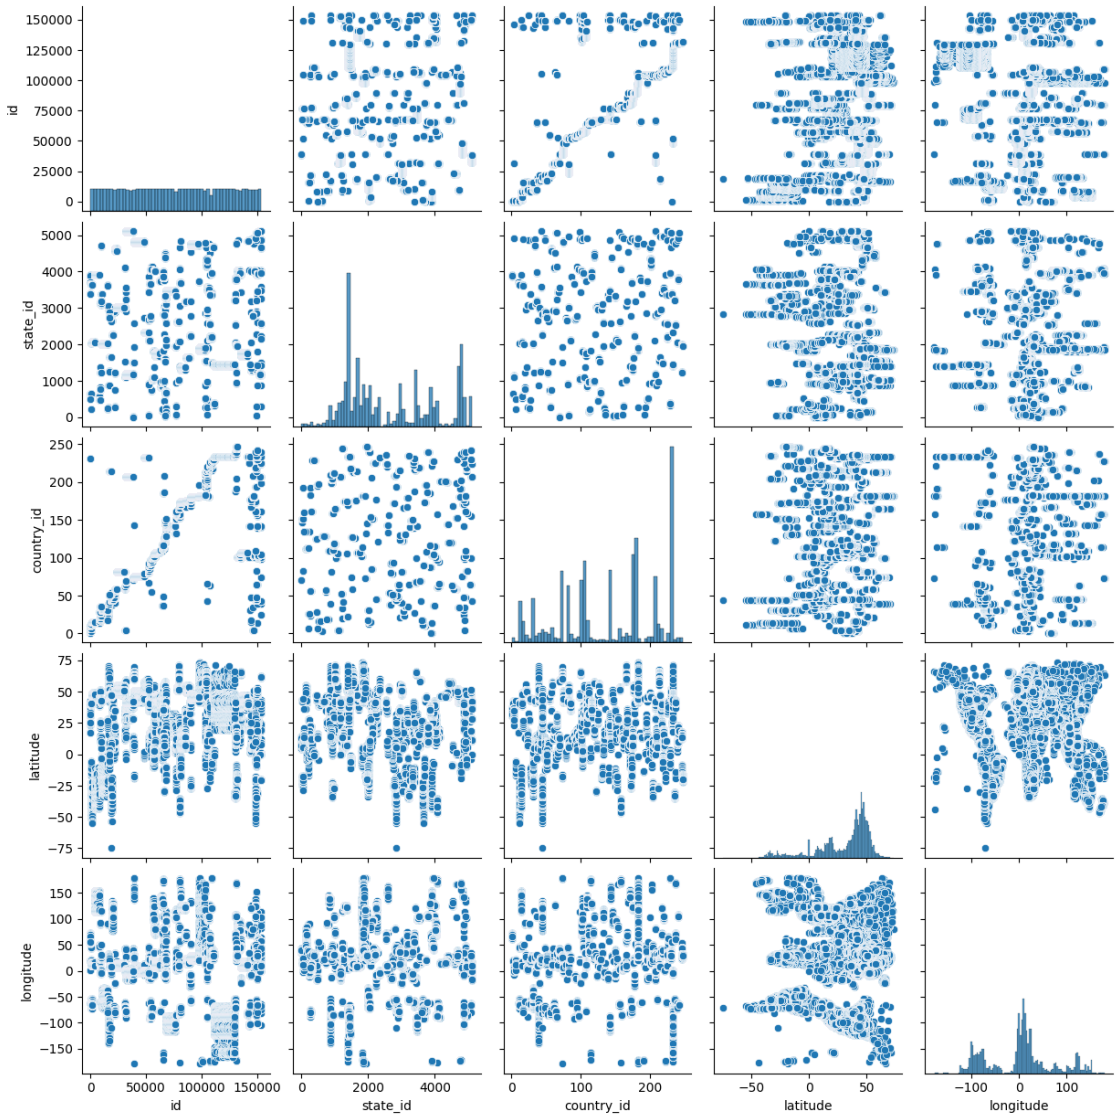

[20]:
```
sns.distplot(df['longitude'])
```

```
<ipython-input-20-4c5c6f107715>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751

  sns.distplot(df['longitude'])
```

[20]: <Axes: xlabel='longitude', ylabel='Density'>

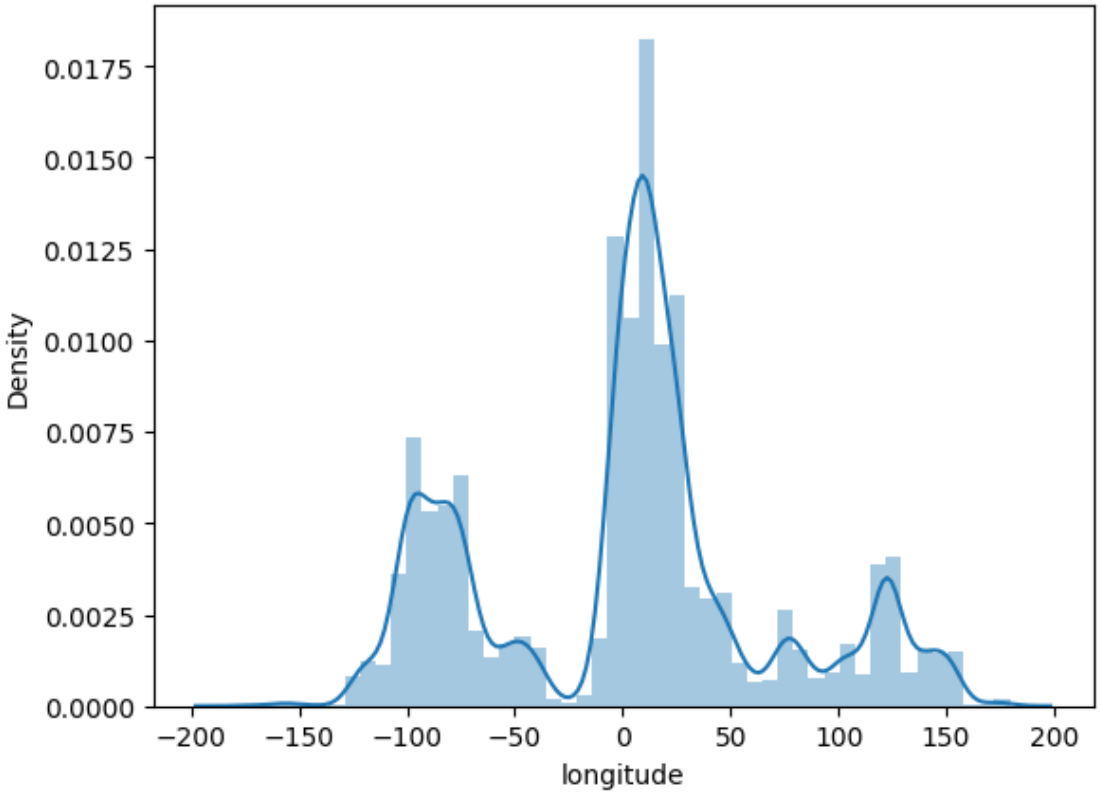

[21]:
```
df1=df[['id', 'state_id', 'country_id',
       'latitude', 'longitude']].dropna()
df1
```

[21]:

	id	state_id	country_id	latitude	longitude
0	52	3901	1	36.68333	71.53333
1	68	3901	1	37.11664	70.58002
2	78	3901	1	36.86477	70.83421
3	84	3901	1	36.95127	72.31800
4	115	3901	1	37.66079	70.67346
…	…	…	…	…	…
150449	131496	1957	247	-19.03333	29.78333
150450	131502	1957	247	-19.78333	29.36667
150451	131503	1957	247	-19.67016	30.00589
150452	131504	1957	247	-19.75000	30.16667
150453	131508	1957	247	-20.30345	30.07514

[150454 rows x 5 columns]

[22]:
```
sns.heatmap(df1.corr())
```

[22]: <Axes: >

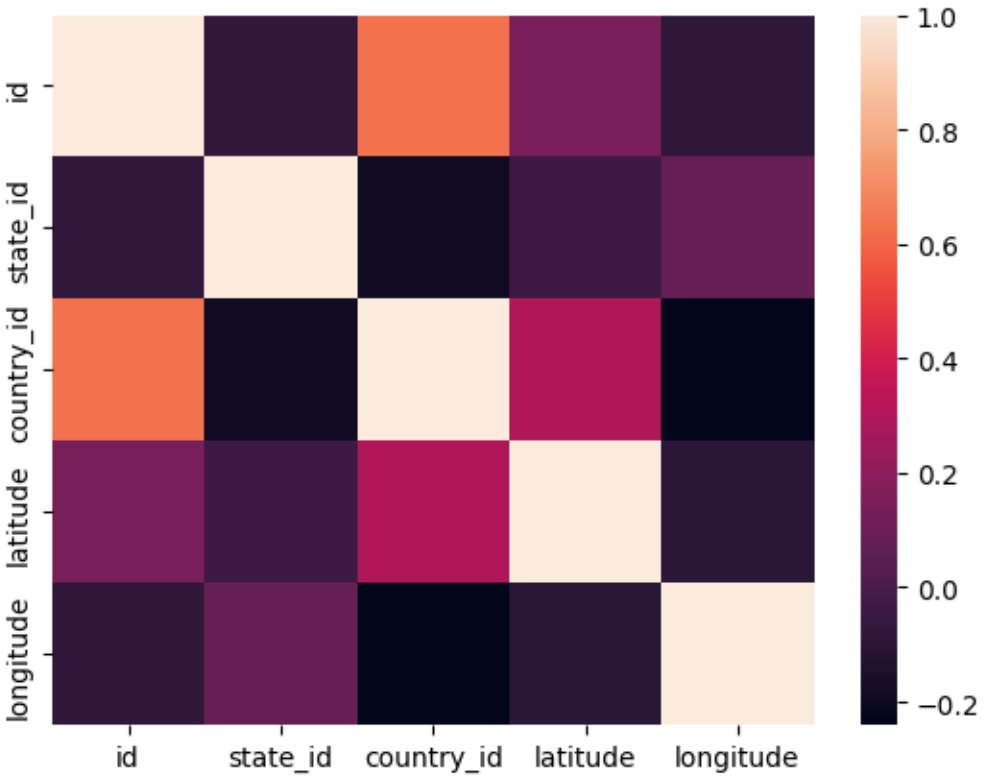

[23]:
```
x=df1[['id', 'state_id', 'country_id',
       'latitude']]
y=df1['longitude']
```

[24]:
```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

[25]:
```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

[25]: LinearRegression()

[26]:
```
print(lr.intercept_)
```

```
26.81844254985726
```

[27]:
```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

[27]:

	Co-efficient
id	0.000154
state_id	0.002057
country_id	-0.276313
latitude	-0.093040

[28]:
```
prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

[28]: <matplotlib.collections.PathCollection at 0x7eea1d2451e0>

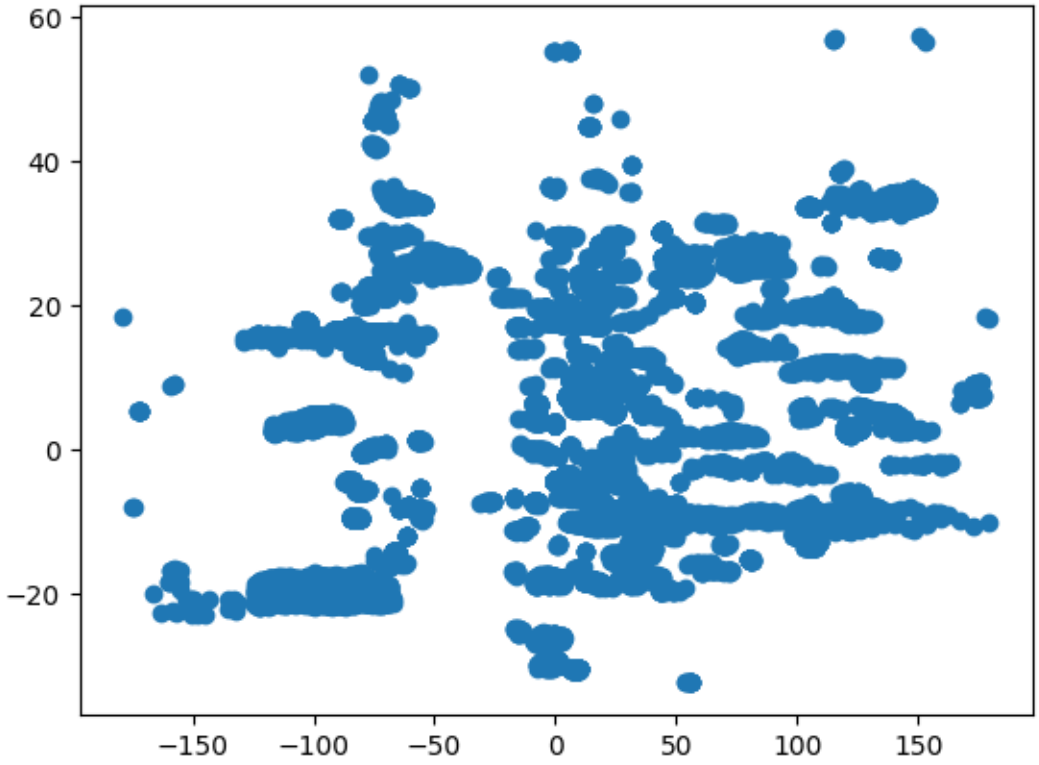

[29]:
```
lr.score(x_test,y_test)
```

[29]: 0.06621717027534668

[30]:
```
lr.score(x_train,y_train)
```

[30]: 0.06723235555568874

[31]:
```
from sklearn.linear_model import Ridge,Lasso
```

[32]:
```
rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

[32]:
```
Ridge(alpha=10)
```

[33]:
```
rr.score(x_test,y_test)
```

[33]:
```
0.06621717026552343
```

[34]:
```
rr.score(x_train,y_train)
```

[34]:
```
0.06723235555568852
```

[35]:
```
la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

[35]:
```
Lasso(alpha=10)
```

[36]:
```
la.score(x_test,y_test)
```

[36]:
```
0.06614419137530625
```

[37]:
```
la.score(x_train,y_train)
```

[37]:
```
0.06718884687477733
```

[38]:
```
from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

[38]:
```
ElasticNet()
```

[39]:
```
en.coef_
```

[39]:
```
array([ 1.53903632e-04,  2.05763581e-03, -2.76206870e-01, -9.20347347e-02])
```

[40]:
```
en.intercept_
```

[40]:
```
26.782642464256952
```

[41]:
```
prediction = en.predict(x_test)
prediction
```

[41]:
```
array([ 25.71503816,   9.90557121,  26.88287221, …,  13.69451066,
       -29.83450702,  15.55788724])
```

[42]:
```
en.score(x_test,y_test)
```

[42]:
```
0.06621553123514157
```

[43]:
```
from sklearn import metrics
```

[44]:
```
print("Mean Absolute Error: ", metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error:  51.58112642435594
```

[45]:
```
print("Mean Squared Error: ", metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error:  4321.822112329533
```

[46]:
```
print("Root Mean Squared Error: ", np.sqrt(metrics.
  ↪mean_squared_error(y_test,prediction)))
```

```
Root Mean Squared Error:  65.74056671743509
```

[47]:
```
import pickle
filename='prediction'
pickle.dump(lr,open(filename,'wb'))
```

[49]:
```
model = pickle.load(open(filename, 'rb'))
real=[[10,20,1,20],[11,23,66,2]]
result = model.predict(real)
result
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does
not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

[49]:
```
array([24.72400315,  8.44468057])
```