

8fg0kzm5z

July 31, 2023

```
[7]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[8]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[9]: df=pd.read_csv("/content/drive/MyDrive/mydatasets/22_countries.csv")
df
```

```
[9]:
```

	id		name	iso3	iso2	numeric_code	phone_code	\
0	1		Afghanistan	AFG	AF	4	93	
1	2		Aland Islands	ALA	AX	248	+358-18	
2	3		Albania	ALB	AL	8	355	
3	4		Algeria	DZA	DZ	12	213	
4	5		American Samoa	ASM	AS	16	+1-684	
..	
245	243	Wallis And Futuna Islands	WLF	WF		876	681	
246	244	Western Sahara	ESH	EH		732	212	
247	245	Yemen	YEM	YE		887	967	
248	246	Zambia	ZMB	ZM		894	260	
249	247	Zimbabwe	ZWE	ZW		716	263	

		capital	currency	currency_name	currency_symbol	tld	\
0		Kabul	AFN	Afghan afghani		.af	
1	Mariehamn		EUR	Euro	€	.ax	
2	Tirana		ALL	Albanian lek	Lek	.al	
3	Algiers		DZD	Algerian dinar		.dz	
4	Pago Pago		USD	US Dollar	\$.as	
..	
245	Mata Utu		XPF	CFP franc		.wf	
246	El-Aaiun		MAD	Moroccan Dirham	MAD	.eh	
247	Sanaa		YER	Yemeni rial		.ye	
248	Lusaka		ZMW	Zambian kwacha	ZK	.zm	

```
249 Harare ZWL Zimbabwe Dollar $ .zw
```

```

native region subregion \
0 Asia Southern Asia
1 Åland Europe Northern Europe
2 Shqipëria Europe Southern Europe
3 Africa Northern Africa
4 American Samoa Oceania Polynesia
.. ...
245 Wallis et Futuna Oceania Polynesia
246 Africa Northern Africa
247 Asia Western Asia
248 Zambia Africa Eastern Africa
249 Zimbabwe Africa Eastern Africa
```

```

timezones latitude longitude \
0 [{zoneName:'Asia/Kabul',gmtOffset:16200,gmtOf... 33.000000 65.0
1 [{zoneName:'Europe/Mariehamn',gmtOffset:7200,... 60.116667 19.9
2 [{zoneName:'Europe/Tirane',gmtOffset:3600,gmt... 41.000000 20.0
3 [{zoneName:'Africa/Algiers',gmtOffset:3600,gm... 28.000000 3.0
4 [{zoneName:'Pacific/Pago_Pago',gmtOffset:-396... -14.333333 -170.0
.. ...
245 [{zoneName:'Pacific/Wallis',gmtOffset:43200,g... -13.300000 -176.2
246 [{zoneName:'Africa/El_Aaiun',gmtOffset:3600,g... 24.500000 -13.0
247 [{zoneName:'Asia/Aden',gmtOffset:10800,gmtOff... 15.000000 48.0
248 [{zoneName:'Africa/Lusaka',gmtOffset:7200,gmt... -15.000000 30.0
249 [{zoneName:'Africa/Harare',gmtOffset:7200,gmt... -20.000000 30.0
```

```

emoji emojiU
0 U+1F1E6 U+1F1EB
1 U+1F1E6 U+1F1FD
2 U+1F1E6 U+1F1F1
3 U+1F1E9 U+1F1FF
4 U+1F1E6 U+1F1F8
.. ...
245 U+1F1FC U+1F1EB
246 U+1F1EA U+1F1ED
247 U+1F1FE U+1F1EA
248 U+1F1FF U+1F1F2
249 U+1F1FF U+1F1FC
```

```
[250 rows x 19 columns]
```

```
[10]: df.head()
```

```

[10]: id name iso3 iso2 numeric_code phone_code capital currency \
0 1 Afghanistan AFG AF 4 93 Kabul AFN
```

1	2	Aland Islands	ALA	AX	248	+358-18	Mariehamn	EUR
2	3	Albania	ALB	AL	8	355	Tirana	ALL
3	4	Algeria	DZA	DZ	12	213	Algiers	DZD
4	5	American Samoa	ASM	AS	16	+1-684	Pago Pago	USD

	currency_name	currency_symbol	tld	native	region \
0	Afghan afghani		.af		Asia
1	Euro	€	.ax	Åland	Europe
2	Albanian lek	Lek	.al	Shqipëria	Europe
3	Algerian dinar		.dz		Africa
4	US Dollar	\$.as	American Samoa	Oceania

	subregion	timezones \
0	Southern Asia	[{zoneName:'Asia\/Kabul',gmtOffset:16200,gmtOf...
1	Northern Europe	[{zoneName:'Europe\/Mariehamn',gmtOffset:7200,...
2	Southern Europe	[{zoneName:'Europe\/Tirane',gmtOffset:3600,gmt...
3	Northern Africa	[{zoneName:'Africa\/Algiers',gmtOffset:3600,gm...
4	Polynesia	[{zoneName:'Pacific\/Pago_Pago',gmtOffset:-396...

	latitude	longitude	emoji	emojiU
0	33.000000	65.0	U+1F1E6	U+1F1EB
1	60.116667	19.9	U+1F1E6	U+1F1FD
2	41.000000	20.0	U+1F1E6	U+1F1F1
3	28.000000	3.0	U+1F1E9	U+1F1FF
4	-14.333333	-170.0	U+1F1E6	U+1F1F8

1 Data Cleaning and Data Preprocessing

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 19 columns):
#   Column              Non-Null Count  Dtype
---  -
0   id                   250 non-null   int64
1   name                 250 non-null   object
2   iso3                 250 non-null   object
3   iso2                 249 non-null   object
4   numeric_code         250 non-null   int64
5   phone_code           250 non-null   object
6   capital              245 non-null   object
7   currency             250 non-null   object
8   currency_name        250 non-null   object
9   currency_symbol      250 non-null   object
10  tld                  250 non-null   object
```

```

11 native          249 non-null    object
12 region          248 non-null    object
13 subregion       247 non-null    object
14 timezones       250 non-null    object
15 latitude        250 non-null    float64
16 longitude       250 non-null    float64
17 emoji          250 non-null    object
18 emojiU         250 non-null    object
dtypes: float64(2), int64(2), object(15)
memory usage: 37.2+ KB

```

```
[12]: df.describe()
```

```

[12]:
count    250.000000    250.000000    250.000000    250.000000
mean     125.500000    435.804000    16.402597    13.52387
std       72.312977    254.38354    26.757204    73.45152
min        1.000000         4.00000   -74.650000   -176.20000
25%       63.250000    219.00000     1.000000   -49.75000
50%      125.500000    436.00000    16.083333    17.00000
75%      187.750000    653.50000    39.000000    48.75000
max      250.000000    926.00000    78.000000   178.00000

```

```
[13]: df.columns
```

```

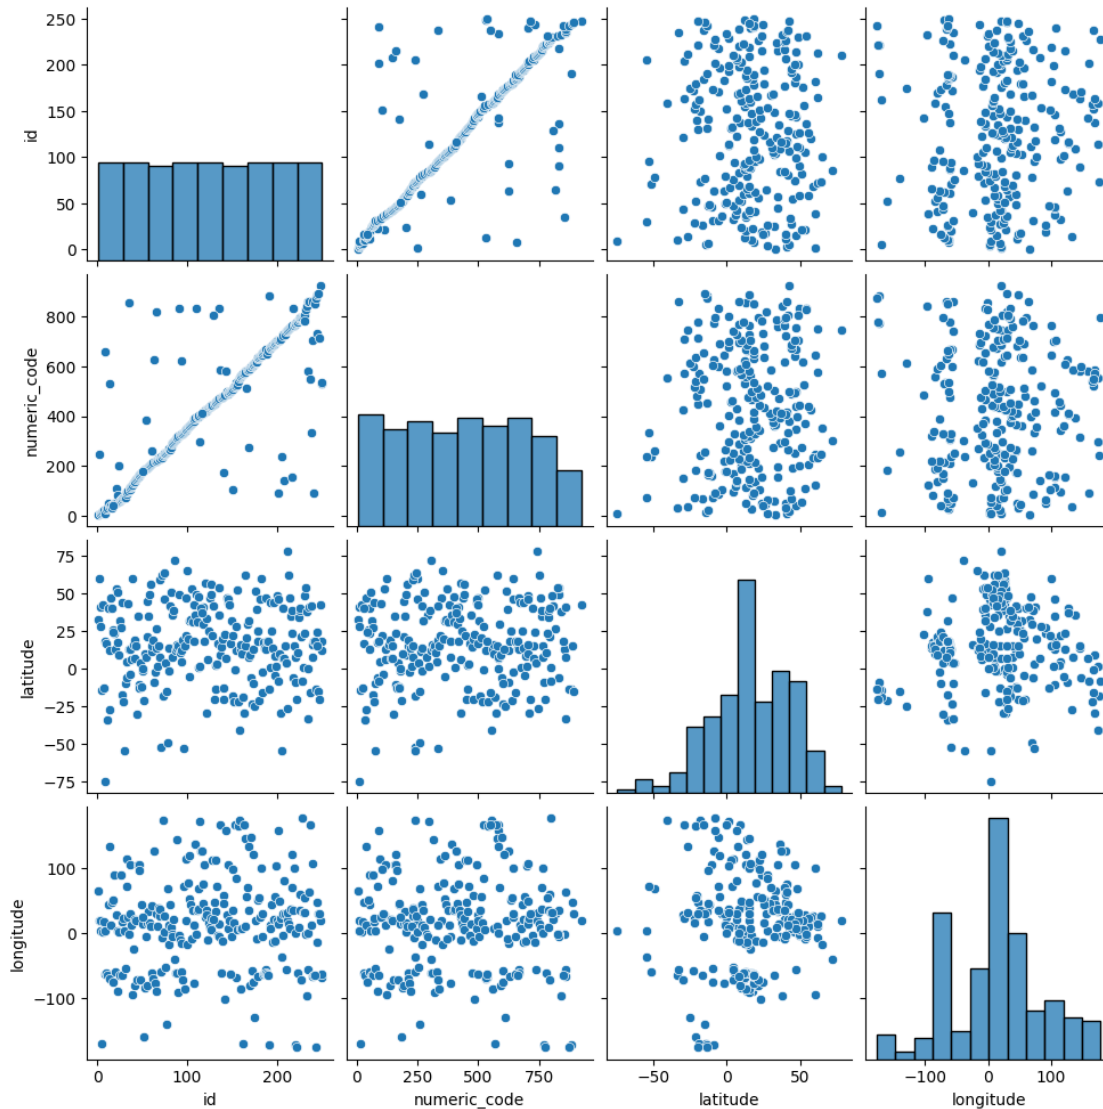
[13]: Index(['id', 'name', 'iso3', 'iso2', 'numeric_code', 'phone_code', 'capital',
            'currency', 'currency_name', 'currency_symbol', 'tld', 'native',
            'region', 'subregion', 'timezones', 'latitude', 'longitude', 'emoji',
            'emojiU'],
            dtype='object')

```

2 EDA and Visualization

```
[14]: sns.pairplot(df)
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x7e48a8969a20>
```



```
[15]: sns.distplot(df['longitude'])
```

<ipython-input-15-4c5c6f107715>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['longitude'])
```

```
[15]: <Axes: xlabel='longitude', ylabel='Density'>
```



```
[16]: df1=df[['id','numeric_code', 'latitude', 'longitude']]
df1
```

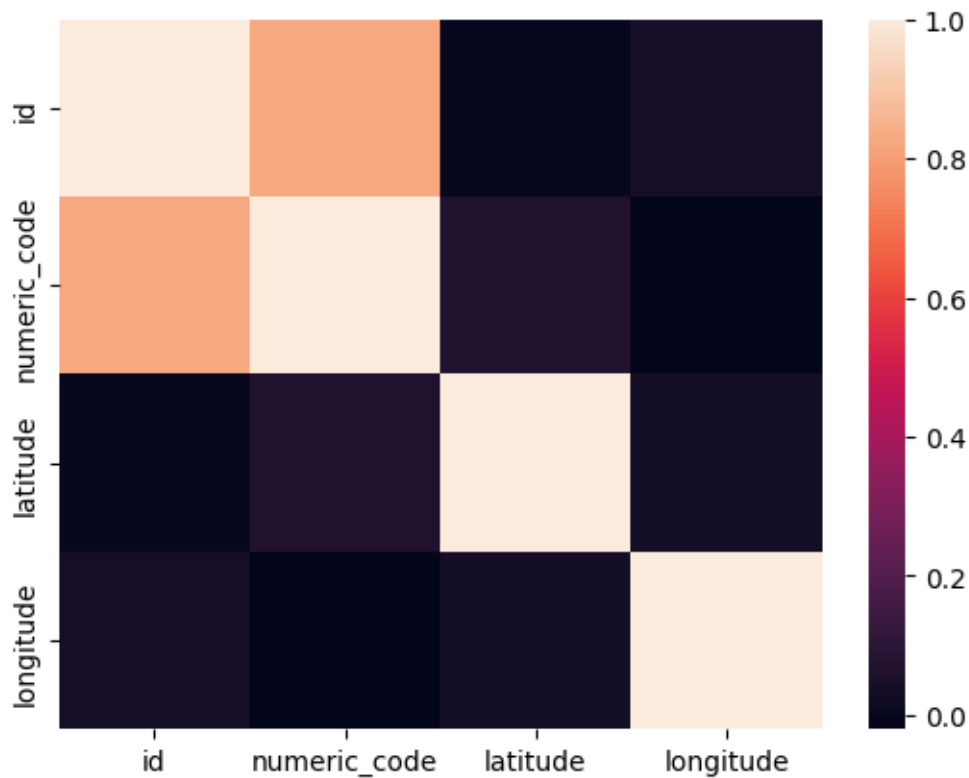
```
[16]:
```

	id	numeric_code	latitude	longitude
0	1	4	33.000000	65.0
1	2	248	60.116667	19.9
2	3	8	41.000000	20.0
3	4	12	28.000000	3.0
4	5	16	-14.333333	-170.0
..
245	243	876	-13.300000	-176.2
246	244	732	24.500000	-13.0
247	245	887	15.000000	48.0
248	246	894	-15.000000	30.0
249	247	716	-20.000000	30.0

[250 rows x 4 columns]

```
[17]: sns.heatmap(df1.corr())
```

[17]: <Axes: >



```
[18]: x=df1[['id','numeric_code', 'latitude']]
      y=df1['longitude']
```

```
[19]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
[20]: from sklearn.linear_model import LinearRegression
      lr=LinearRegression()
      lr.fit(x_train,y_train)
```

[20]: LinearRegression()

```
[21]: print(lr.intercept_)
```


9.24756110032656

```
[22]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
      coeff
```

```
[22]:          Co-efficient
      id          0.120467
      numeric_code -0.031209
      latitude     0.106156
```

```
[23]: prediction =lr.predict(x_test)
      plt.scatter(y_test,prediction)
```

```
[23]: <matplotlib.collections.PathCollection at 0x7e489d28e8f0>
```



```
[24]: lr.score(x_test,y_test)
```

```
[24]: 0.02021428278698778
```

```
[25]: lr.score(x_train,y_train)
```

```
[25]: 0.006070210320557434
```

```
[26]: from sklearn.linear_model import Ridge,Lasso
```

```
[27]: rr=Ridge(alpha=10)
      rr.fit(x_train,y_train)
```


[27]: Ridge(alpha=10)

```
[28]: rr.score(x_test,y_test)
```

[28]: 0.02021367175113542

```
[29]: rr.score(x_train,y_train)
```

[29]: 0.006070210305520685

```
[30]: la=Lasso(alpha=10)
      la.fit(x_train,y_train)
```

[30]: Lasso(alpha=10)

```
[31]: la.score(x_test,y_test)
```

[31]: 0.01910812291362196

```
[32]: la.score(x_train,y_train)
```

[32]: 0.0060280103999195145

```
[33]: from sklearn.linear_model import ElasticNet
      en=ElasticNet()
      en.fit(x_train,y_train)
```

[33]: ElasticNet()

```
[34]: en.coef_
```

[34]: array([0.12008244, -0.03110431, 0.10530327])

```
[35]: en.intercept_
```

[35]: 9.264537380671829

```
[36]: prediction = en.predict(x_test)
      prediction
```

[36]: array([15.89697488, 11.23015295, 12.67089952, 12.24363451, 12.47077855,
 6.66121143, 8.62350019, 12.6164041 , 7.32140493, 15.15398186,
 13.3784481 , 10.93563319, 13.19420189, 9.85967827, 12.42579498,
 10.56643768, 14.7649501 , 9.57013269, 16.00574016, 8.66545413,
 15.23311907, 13.61510864, 12.03214058, 9.5541292 , 12.65113159,
 11.39266148, 9.01010883, 32.76248961, 17.11788031, 14.54814853,
 8.83415584, 12.39541688, 14.78120231, 5.35766089, 15.38881248,

```
15.93618662, 14.63000845, 14.23125461, 15.39393645, 12.05408854,
11.88587358, 8.98416373, 13.68373316, 18.37622796, 13.69338444,
16.6060118 , 13.6984063 , -4.4366965 , 9.45238614, 13.89623932,
16.71456217, 16.53945486, 4.05918149, 14.07396214, 16.36407217,
13.29184498, 9.62678895, 4.81200829, 9.15057233, 8.77224966,
10.16197492, 12.67476147, 10.87658947, 11.54099671, 13.2880713 ,
10.71510886, 12.51816758, 12.37038299, 11.64691845, 14.33755451,
15.88346382, 14.80421094, 15.62848939, 14.2470795 , 11.38986191]]
```

```
[37]: en.score(x_test,y_test)
```

```
[37]: 0.020155445810723704
```

```
[38]: from sklearn import metrics
```

```
[39]: print("Mean Absolute Error: ", metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 50.43961554465261
```

```
[40]: print("Mean Squared Error: ", metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 4380.921206445248
```

```
[41]: print("Root Mean Squared Error: ", np.sqrt(metrics.
↪mean_squared_error(y_test,prediction)))
```

```
Root Mean Squared Error: 66.18852775553516
```

```
[42]: import pickle
filename='prediction'
pickle.dump(lr,open(filename,'wb'))
```

```
[43]: model = pickle.load(open(filename, 'rb'))
real=[[10,20,30],[52,23,66]]
result = model.predict(real)
result
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does
not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

```
[43]: array([13.01274813, 21.80037036])
```